

## Lecture 14: Dijkstra's Algorithm

### Today: SSSP without negative weights

- If all weights are  $\geq 0$ , can do better than Bellman Ford's  $O(|V| \cdot |E|)$
- Dijkstra's Algorithm takes  $O(|E| + |V| \log |V|)$  time

### Idea of Dijkstra's Algorithm

- Intuition: expanding frontier
  - At time  $x$ , identify all nodes with  $\delta \leq x$
  - Frontier takes  $w(u, v)$  time to traverse edge  $(u, v)$
  - Grow frontier until all nodes are reached
- Familiar case:
  - If all weights are 1, this is BFS
  - For small integer weights, can subdivide edges and then use BFS
    - \* like Wleft Brothers in pset problem 6-4(a)
    - \* Positive integer weights  $\leq k$  leads to  $O(|V| + k \cdot |E|)$  runtime.
    - \* Bad if weights are big; doesn't work if weights are fractional
- Goal: Simulate continuous expansion process with discrete steps
  - Fast-forward to the next **event**: next time the frontier reaches a new vertex
  - Priority queue!
- Example (see slides)

### Dijkstra's Algorithm

- Fits in relaxation framework
- Uses a Priority Queue to decide relaxation order
- Need a changeable Priority Queue! How?
  - Pset problem 5-2 showed how to cross-link a binary heap with a DAA (or Hash Table) to support `change_key` in  $O(\log n)$  time (expected time if using a Hash Table)

- \* Recall: `Q.change_key(id, new_key)` finds the item with identifier `id`, changes its key to `new_key`, and readjusts the queue as necessary
- \* `decrease_key` is the special case where the new key must be smaller than the old key
- Easier DAA-based priority queue can easily be modified in  $O(1)$  time

```

1 def dijkstra_sssp(Adj, w, s):
2     parent = [None] * len(Adj) # Same
3     parent[s] = s               # init
4     d = [math.inf] * len(Adj)  # as
5     d[s] = 0                   # before.
6
7     Q = PriorityQueue.build(Item(id=u, key=d[u]) for u in Adj)
8
9     while len(Q) > 0:
10        u = Q.delete_min().id # Delete and process u
11        for v in Adj[u]:      # Same
12            if d[v] > d[u] + w(u,v): # relax
13                d[v] = d[u] + w(u,v) # as
14                parent[v] = u         # before.
15                Q.decrease_key(id=v, new_key=d[v]) # NEW!
16
17    return d, parent

```

- Runtime is dominated by Priority Queue operations: build once, `delete_min`  $|V|$  times, and `decrease_key`  $|E|$  times.
- Which Priority Queue to use?! Four standard contenders (all runtimes are  $O(\dots)$ ):

Priority Queue	build	del_min	dec_key	Total
DAA Heap	$V$	$V$	1	$V^2$
Hash Table Heap	$V$ ex.	$V$ ex.	1 ex.	$V^2$ ex.
Binary Heap	$V$	$\log V$	$\log V$	$(V + E) \log V$
Fibonacci Heap	$V$	$\log V$ am.	1 am.	$E + V \log V$

- This class **does not cover** how Fibonacci Heaps work; you only need to know these runtimes
- When to use each?
  - Fibonacci Heap: Asymptotically best of the three options, but higher constant factors
  - Hash Table / Direct Access Array: Tied with Fibonacci Heaps for dense graphs with  $E = \Omega(V^2)$ .
  - Binary Heap: Tied with Fibonacci Heaps for sparse graphs with  $E = O(V)$ , e.g., planar graphs.
  - In the real world, Binary Heaps are often faster<sup>[citation needed]</sup>, even if not asymptotically

## Correctness of Dijkstra

- Intuition for the proof: the frontier includes nodes with “small”  $\delta$  and excludes nodes with “large”  $\delta$ . How to formalize this?
- Let  $B$  be the set of vertices that have been deleted from the Queue. Every loop adds one node to  $B$ .
- Loop Invariant:
  - For each node  $u \in B$ ,  $d[u] = \delta(s, u)$ .
  - For each node  $v \notin B$ , its key  $d[v]$  in  $Q$  equals the minimum weight of an  $s$ - $v$  path that **only visits nodes in  $B$**  except for  $v$
- Base Case: True at beginning when  $B = \emptyset$ , since there's only one (zero-edge) path that stays in  $B$  until the last vertex, and that path has weight 0.
- Inductive step: assume  $B$  is not empty and satisfies our loop invariant. Let  $M = \min\{d[v] \mid v \notin B\}$  be the minimum key in  $Q$ . (Note: Since keys are deleted in increasing order, we have  $\delta(s, u) \leq M$  for every  $u \in B$ .)
- Claim: All paths from  $s$  that leave  $B$  have weight  $\geq M$ . Proof: Let  $x$  be the first node the path visits outside of  $B$ . Then the path has weight  $\geq d[x] \geq M$ .
- All nodes  $v \notin B$  have  $\delta(s, v) \geq M$ . Proof: any  $s$ - $v$  path leaves  $B$ .
- The next vertex  $u_{\text{next}}$  that is about to get deleted from  $Q$  has  $d[u_{\text{next}}] = M$ , since that is `delete_min`'s job.
- We just argued that  $\delta(s, u_{\text{next}}) \geq M$ , but by safety  $\delta(s, u_{\text{next}}) \leq d[u_{\text{next}}] = M$ , so indeed  $d[u_{\text{next}}] = \delta(s, u_{\text{next}}) = M$ .
- What about the rest of the nodes  $v \notin B$  and  $v \neq u_{\text{next}}$ ? What can the shortest  $s$ - $v$  path that doesn't leave  $B \cup \{u_{\text{next}}\}$  look like?
  - Case 1: the penultimate vertex is  $u_{\text{next}}$ . The best weight of such a path is  $\delta(s, u_{\text{next}}) + w(u_{\text{next}}, v)$ .
  - Case 2: the penultimate vertex is some  $u \neq u_{\text{next}}$ . Then replace the path to  $u$  with a shortest path that stays entirely in  $B$ , so this new path to  $v$  doesn't visit  $u_{\text{next}}$  at all. The smallest weight of such a path is (the old value of)  $d[v]$ .
- Since edge  $(u_{\text{next}}, v)$  gets relaxed when  $u_{\text{next}}$  is processed, the loop invariant is restored.

**Single Pair Search: Compute a single  $\delta(s, t)$** 

- What if we just need  $\delta(s, t)$ , don't care about the rest of the graph
- Idea: Search from both directions at once! When frontiers overlap, we've found the shortest path.
- Hopefully we've avoided looking at much of the graph.
- In more detail:
  - Run two Dijkstra's: one from  $s$ , and one backwards from  $t$  (with all edges reversed).
  - Two separate Priority Queues  $Q_s$  and  $Q_t$ , two parent pointers  $\text{parent}_s[v]$  and  $\text{parent}_t[v]$ , and two distance measures  $d_s[v]$  and  $d_t[v]$
  - Interleave steps from the two Dijkstra runs, in any order
  - Stop after some vertex has been deleted and processed in both runs
  - **Warning:** This first vertex to be processed twice **might not** be on the shortest path!
  - Read through all processed vertices and find the vertex  $x$  that minimizes  $d_s[x] + d_t[x]$   
Use parent pointers to reconstruct and return the  $s$ - $t$  path through  $x$ .
- Runtime: No guaranteed asymptotic gains, but often great in practice

SSSP Algorithm	Lecture	Setting	Running Time
BFS	10	$w(e) = 1$ for all $e \in E$	$O(V + E)$
DAG shortest paths	12	$G$ acyclic	$O(V + E)$
Dijkstra	14	$w(e) \geq 0$ for all $e \in E$	$O(V \log V + E)$
Bellman-Ford	13	(general)	$O(VE)$

```
def dijkstra_sssp(Adj, w, s):
    parent = [None] * len(Adj)    # Same
    parent[s] = s                  # init
    d = [math.inf] * len(Adj)     # as
    d[s] = 0                       # before.

    Q = PriorityQueue.build(Item(id=u, key=d[u]) for u in Adj)

    while len(Q) > 0:
        u = Q.delete_min().id    # Delete and process u
        for v in Adj[u]:         # Same
            if d[v] > d[u] + w(u,v): # relax
                d[v] = d[u] + w(u,v) # as
                parent[v] = u         # before.
                Q.decrease_key(id=v, new_key=d[v]) # NEW!

    return d, parent
```

# Dijkstra Runtime

Build +  $V \cdot \text{DeleteMin}$  +  $E \cdot \text{DecreaseKey}$

Queue	Build	Delete Min	Decrease Key	Total
DAA Heap	$V$	$V$	1	$V^2 + E = O(V^2)$
Hash Table Heap	$V$ ex.	$V$ ex.	1 ex.	$V^2$ ex.
Bin Heap	$V$	$\log V$	$\log V$	$(V + E) \log V$
<i>Fibonacci Heap</i>	$V$	$\log V$ am.	<b>1 am.</b>	$V \log V + E$

(All runtimes are big-O)



# Dijkstra vs Bidirectional Dijkstra







